

Simple Data Compression

CS 10C Programming Assignment

Huffman coding is used to compress data. The idea is straightforward: represent more common longer strings with shorter ones via a basic translation matrix. The translation matrix is easily computed from the data itself by counting and sorting by frequency.

For example, in a well-known corpus used in Natural Language Processing called the "Brown" corpus (see nltk.org), the top-20 most frequent tokens, which are words or punctuation marks are listed below associated with frequency and code. The word "and" for example requires writing three characters. However, if I encoded it differently, say, using the word "5" (yes, I called "5" a word on purpose), then I save having to write two extra characters! Note, the word "and" is so frequent, I save those two extra characters many times over!

Token	Frequency	Code
the	62713	1
,	58334	2
.	49346	3
of	36080	4
and	27932	5
to	25732	6
a	21881	7
in	19536	8
that	10237	9
is	10011	10
was	9777	11
for	8841	12
``	8837	13
"	8789	14
The	7258	15
with	7012	16
it	6723	17
as	6706	18
he	6566	19
his	6466	20

So the steps of Huffman coding are relatively straightforward:

1. Pass through the data once, collecting a list of token-frequency counts.
2. Sort the token-frequency counts by frequency, in descending order.
3. Assign codes to tokens using a simple counter, for example by incrementing over the integers; this is just to keep things simple.

4. Store the new mapping (token -> code) in a hashtable called "encoder".
5. Store the reverse mapping (code -> token) in a hashtable called "decoder".
6. Pass through the data a second time. This time, replace all tokens with their codes.

Now, be amazed at how much you've shrunk your data!

Delivery Notes:

- (1) Implement your own hashtable from scratch, you are not allowed to use existing hash table libraries.
- (2) To be useful, your output should include the coded data as well as the decoder (code -> token) mapping file.

Now GZIP all that and watch it shrink immensely!